

# AI 用于软件安全和漏洞挖掘

陈浩

University of California, Davis, USA

2020-12-04

# Success of AI

- AI has seen great success in
  - vision
  - speech
  - NLP
- How useful is AI for security?

# Software security

## Traditional methods

- Static analysis
- Dynamic analysis
- Symbolic execution
- Model checking

Can we apply AI to software security?

# Fuzzing

- A popular dynamic analysis technique
- Mutates inputs to trigger bugs in the program.
- Benefits
  - No false positive.
  - Provides witness inputs to trigger bugs.
  - Scalable
- Challenge: How to explore different states of the program?

# Background

## American Fuzzy Lop (AFL)

- Instruments the program to track branch coverage.
- Keeps inputs that explore new branches as seeds.
- Mutates inputs by heuristics.

## Limitations

- Uses only heuristics, with no principled guidance.
- Cannot target specific branches.
- Random mutation of input is unproductive.

# Fuzzing: from art to science

- Model fuzzing as an optimization problem.
- Apply principled mathematical tools.

## Motivating example

```
void foo(int i, int j) {  
  
    if (i * i - j * 2 > 0) {  
        // some code  
    } else {  
        // some code  
    }  
}  
  
int main() {  
    char buf[1024];  
    int i = 0, j = 0;  
    if (fread(buf, sizeof(char), 1024, fp) < 1024) {  
        return(1);  
    }  
    if (fread(&i, sizeof(int), 1, fp) < 1) {  
        return(1);  
    }  
    if (fread(&j, sizeof(int), 1, fp) < 1) {  
        return(1);  
    }  
    foo(i, j);  
}
```

## Motivating example

```
void foo(int i, int j) {  
    // Byte-level taint tracking: which input bytes flow into this predicate?  
    if (i * i - j * 2 > 0) {  
        // some code  
    } else {  
        // some code  
    }  
}  
  
int main() {  
    char buf[1024];  
    int i = 0, j = 0;  
    if (fread(buf, sizeof(char), 1024, fp) < 1024) {  
        return(1);  
    }  
    if (fread(&i, sizeof(int), 1, fp) < 1) {  
        return(1);  
    }  
    if (fread(&j, sizeof(int), 1, fp) < 1) {  
        return(1);  
    }  
    foo(i, j);  
}
```



## Motivating example

```
void foo(int i, int j) {  
    // Search based on gradient descent: how to solve this path constraint efficiently?  
    if (i * i - j * 2 > 0) {  
        // some code  
    } else {  
        // some code  
    }  
}  
  
int main() {  
    char buf[1024];  
    int i = 0, j = 0;  
    if (fread(buf, sizeof(char), 1024, fp) < 1024) {  
        return(1);  
    }  
    if (fread(&i, sizeof(int), 1, fp) < 1) {  
        return(1);  
    }  
    if (fread(&j, sizeof(int), 1, fp) < 1) {  
        return(1);  
    }  
    foo(i, j);  
}
```

# Byte-level taint tracking

**Question** Which input bytes flow into each path constraint?

**Challenges** Taint tracking is expensive, more so at byte-level.

**Observation** Taint tracking is unnecessary while mutating the input.

- Solution**
- Run taint tracking **once** on each seed input.
  - Mutate the seed **many times** and run **without** taint tracking.
  - Benefit: amortize the cost of taint tracking over many mutations.

# Search based on gradient descent

**Question** How to mutate the input to solve path constraints?

- Solution**
- View each path constraint as a function over the relevant input bytes.
  - Solve the constraint using optimization techniques, such as gradient descent.

# View path constraint as function over input


View each path constraint as  $f(x)$  where

- $x$ : a vector representing the relevant input bytes
- $f(x)$ 
  - Represents the computation from the program start.
  - Needs no analytic form.

Comparison	$f(\cdot)$	Constraint
$a < b$	$a - b$	$f(\cdot) < 0$
$a \leq b$	$a - b$	$f(\cdot) \leq 0$
$a > b$	$b - a$	$f(\cdot) < 0$
$a \geq b$	$b - a$	$f(\cdot) \leq 0$
$a = b$	$\text{abs}(a - b)$	$f(\cdot) = 0$
$a \neq b$	$-\text{abs}(a - b)$	$f(\cdot) < 0$

# Gradient descent

- Goal: find a minimum of  $f(x)$ .
- Iterative algorithm
  1. Start:  $x \leftarrow x_0$
  2. Repeat
    - 2.1 Compute  $\nabla_x f(x)$
    - 2.2  $x \leftarrow x - \epsilon \nabla_x f(x)$  where  $\epsilon$ : learning rate



# Apply gradient descent to solving path constraints

## Advantages

- In machine learning, the goal is to find the global minimum.
- In fuzzing, need only find a good enough solution  $f(x) < 0$ .

## Challenges

In fuzzing, we cannot compute gradient directly because

- $f(x)$  has no closed form.
- $f(x)$  is not continuous because  $x$  is usually discrete.

# Numerically approximate directional derivative

## Directional derivative

$$\frac{\partial f(x)}{\partial x_i} = \lim_{\delta \rightarrow 0} \frac{f(x + \delta e_i) - f(x)}{\delta}$$

## Approximation

Let  $\delta$  be the smallest values of the type (usually 1 or -1).

## Algorithm

1. Run program with the input  $x$  to get  $f(x)$
2. For  $i \in [1, n]$ 
  - 2.1 Run program with the input  $x + \delta e_i$  to get  $f(x + \delta e_i)$

2.2

$$\frac{\partial f(x)}{\partial x_i} \approx \frac{f(x + \delta e_i) - f(x)}{\delta}$$

## Nested conditional statements

```
// pngutil.c 2406
png_crc_read(png_ptr, buffer, length);
buffer[length] = 0;

if (png_crc_finish(png_ptr, 0) != 0)
    return;

if (buffer[0] != 1 && buffer[0] != 2)
{
    png_chunk_benign_error(png_ptr, "invalid unit");
    return;
}
```



## Nested conditional statements

```
// pngutil.c 2406
png_crc_read(png_ptr, buffer, length);
buffer[length] = 0;
// First conditional statement
if (png_crc_finish(png_ptr, 0) != 0)
    return;
// Second conditional statement
if (buffer[0] != 1 && buffer[0] != 2)
{
    png_chunk_benign_error(png_ptr, "invalid unit");
    return;
}
```

## Nested constraints are difficult to solve

Program	Percentage of nested constraints in	
	all unsolved constraints	all constraints
<i>djpeg</i>	90.00 %	75.65 %
<i>file</i>	86.49 %	44.14 %
<i>jhead</i>	57.95 %	51.53 %
<i>mutool</i>	80.88 %	58.63 %
<i>nm</i>	84.32 %	68.16 %
<i>objdump</i>	90.54 %	73.95 %
<i>readelf</i>	84.12 %	70.50 %
<i>readpng</i>	94.02 %	89.50 %
<i>size</i>	87.86 %	71.46 %
<i>tcpdump</i>	96.15 %	78.98 %
<i>tiff2ps</i>	75.56 %	62.18 %
<i>xmlint</i>	78.18 %	72.37 %
<i>xmlwf</i>	96.18 %	68.16 %

# Control and data flow dependencies of nested conditional statements

```
void foo(unsigned x, unsigned y, unsigned z) {  
  
    if (w + x < 2) {  
  
        if (x + y < 3) {  
  
            if (a == 1111) {  
                if (y + z == 2222) {  
                    ...  
                }  
                if (y > 1) {  
                    ...  
                } else {  
                    ...  
                }  
            }  
        }  
    }  
}
```

# Control and data flow dependencies of nested conditional statements

```
void foo(unsigned x, unsigned y, unsigned z) {  
    // control flow dependency  
    if (w + x < 2) {  
        // control flow dependency  
        if (x + y < 3) {  
            // control flow dependency  
            if (a == 1111) {  
                if (y + z == 2222) {  
                    ...  
                }  
                if (y > 1) {  
                    ...  
                } else {  
                    ...  
                }  
            }  
        }  
    }  
}
```

# Control and data flow dependencies of nested conditional statements

```
void foo(unsigned x, unsigned y, unsigned z) {  
    // data flow dependency  
    if (w + x < 2) {  
        // data flow dependency  
        if (x + y < 3) {  
            if (a == 1111) {  
                if (y + z == 2222) {  
                    ...  
                }  
                if (y > 1) {  
                    ...  
                } else {  
                    ...  
                }  
            }  
        }  
    }  
}
```

# Solve nested constraints

1. Determine control flow dependency
  - based on post-dominator
2. Determine data flow dependency
  - based on dynamic taint analysis
3. Solve constraints
  - Prioritize reachability
  - Prioritize satisfiability
  - Joint optimization

# Joint optimization

- Problem: find  $x$  such that  $f_i(x) = 0, \forall i \in [1, n]$
- Define

$$g(x) = \sum_{i=1}^n R(f_i(x))$$

where the rectifier  $R(x) \equiv 0 \vee x$

- $g(x) = 0 \Rightarrow f_i(x) \leq 0, \forall i \in [1, n]$
- Use gradient descent to solve the constraint.

# Joint optimization

```
void foo(unsigned x, unsigned y, unsigned z) {  
  
    if (w + x < 2) {  
  
        if (x + y < 3) {  
            if (a == 1111) {  
  
                if (y + z == 2222) {  
                    ...  
                }  
            }  
        }  
    }  
}
```



# Joint optimization

```
void foo(unsigned x, unsigned y, unsigned z) {  
    //  $f_1(\cdot) = w + x - 2 + \epsilon$   
    if (w + x < 2) {  
        //  $f_2(\cdot) = x + y - 3 + \epsilon$   
        if (x + y < 3) {  
            if (a == 1111) {  
                //  $f_3(\cdot) = y + z - 2222$   
                if (y + z == 2222) {  
                    ...  
                }  
            }  
        }  
    }  
}
```

## Joint optimization

```
void foo(unsigned x, unsigned y, unsigned z) {  
    //  $f_1(\cdot) = w + x - 2 + \epsilon$   
    if (w + x < 2) {  
        //  $f_2(\cdot) = x + y - 3 + \epsilon$   
        if (x + y < 3) {  
            if (a == 1111) {  
                //  $f_3(\cdot) = y + z - 2222$   
                if (y + z == 2222) {  
                    ...  
                }  
            }  
        }  
    }  
}
```

$$g(\cdot) \equiv \sum_{i=1}^3 R(f_i(\cdot))$$

$g(\cdot) = 0 \Rightarrow$  all constraints are satisfied

# Implementation

## Instrumentation

- Instrument LLVM Pass.
- Extend DFSAN with byte-level taint tracking.
- Size: 1262 line C++

## Fuzzer

Size: 8672 line Rust

## Bugs found on the LAVA-M data set

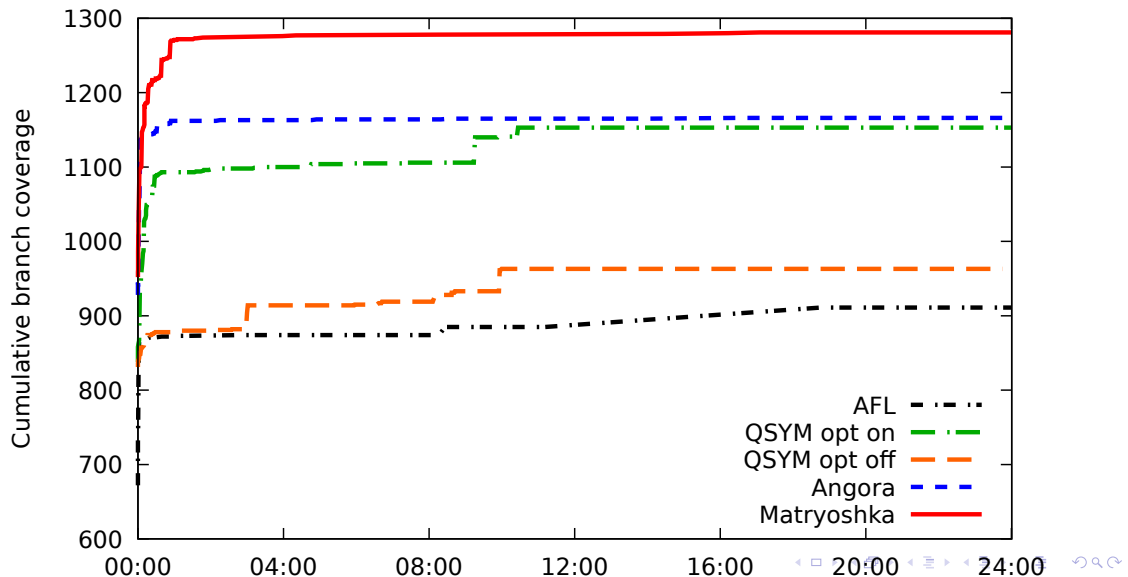
Program	Listed bugs	Bugs found by					
		AFL	QSYM	NEUZZ	REDQUEEN	Angora	<b>Matryoshka</b>
<i>uniq</i>	28	9	28	29	29	29	29
<i>base64</i>	44	0	44	48	48	48	48
<i>md5sum</i>	57	0	57	60	57	57	57
<i>who</i>	2136	1	1238	1582	2462	1541	2432

## Verified bugs

Program	Types of bugs				Total
	SBO	HBO	OOM	OBR	
<i>file</i>	4				4
<i>jhead</i>	2	15		6	23
<i>nm</i>	1	1			2
<i>objdump</i>			3	1	4
<i>size</i>		1	1		2
<i>readelf</i>		4			4
<i>tiff2ps</i>		1	1		2

- SBO: stack buffer overflow
- HBO: heap buffer overflow
- OOM: out of memory
- OBR: out of bound read

## Cumulative branch coverage on *readpng*



# Conclusion

- AI is a powerful tool for software security.
- Formulate security problem as AI/optimization problem.
- AI and traditional program analyses are mutually beneficial.

Angora



Matryoshka

